

# M4 – Parallelism

Implementation of Locks  
Cache Coherence

# Outline

- Parallelism
- Flynn's classification
- Vector Processing
  - Subword Parallelism
- Symmetric Multiprocessors, Distributed Memory Machines
  - Shared Memory Multiprocessing, Message Passing
- Synchronization Primitives
  - Locks, LL-SC
- Cache coherence

# Synchronization

- Two processors sharing an area of memory
  - P1 writes, then P2 reads
  - Data race if P1 and P2 don't synchronize
    - Result depends of order of accesses

# Synchronization

- Two processors sharing an area of memory
  - P1 writes, then P2 reads
  - Data race if P1 and P2 don't synchronize
    - Result depends of order of accesses
- Hardware support required
  - Atomic read/write memory operation
  - No other access to the location allowed between the read and write

# Synchronization

- Two processors sharing an area of memory
  - P1 writes, then P2 reads
  - Data race if P1 and P2 don't synchronize
    - Result depends of order of accesses
- Hardware support required
  - Atomic read/write memory operation
  - No other access to the location allowed between the read and write
- Could be a single instruction
  - E.g., atomic swap of register ↔ memory
  - Or an atomic pair of instructions

ISA

lock L

# Implementing Locks

- Must synchronize processes so that they access shared variable one at a time in critical section; called **Mutual Exclusion**

# Implementing Locks

- Must synchronize processes so that they access shared variable one at a time in critical section; called **Mutual Exclusion**
- Mutex Lock: a synchronization primitive

# Implementing Locks

- Mutex Lock: a synchronization primitive

- AcquireLock(L)

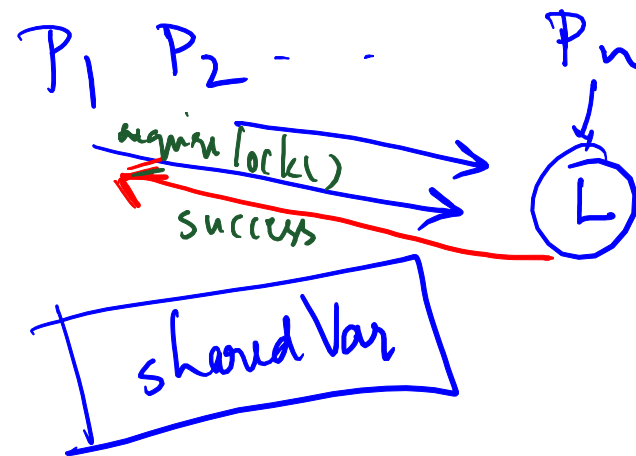
–

–

- ReleaseLock(L)

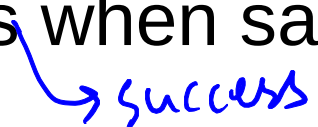
–

–

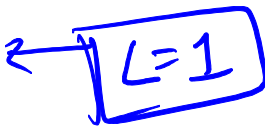


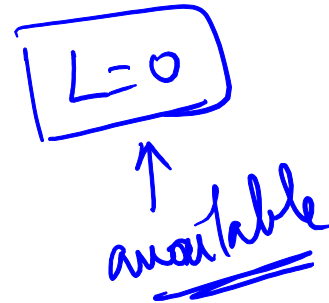



# Implementing Locks

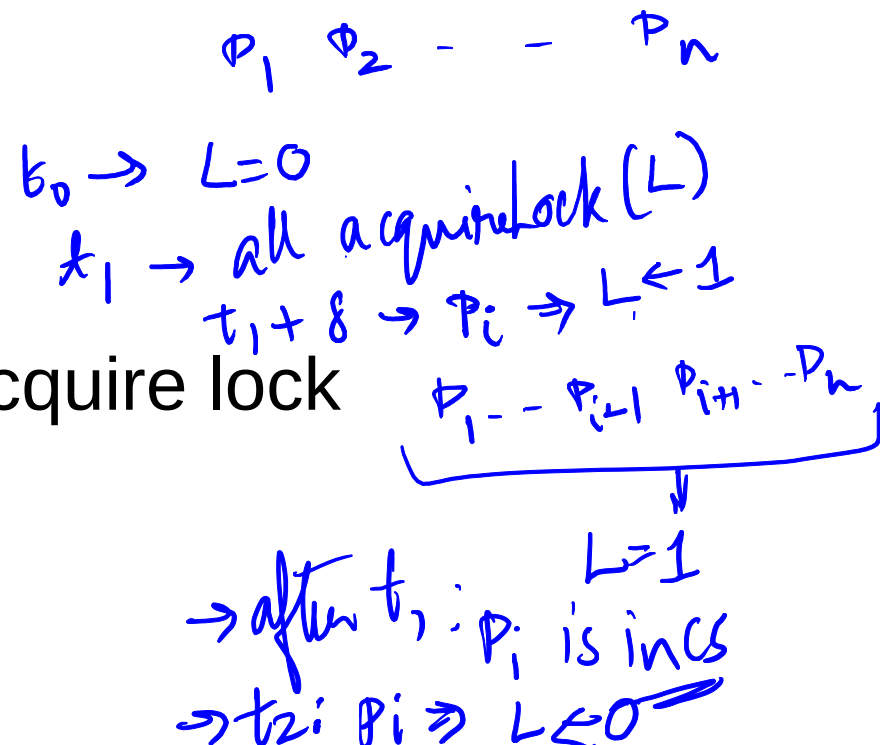
- Mutex Lock: a synchronization primitive
- AcquireLock(L)
  - Done before critical section of code
  - Returns when safe for process to enter critical section 
- ReleaseLock(L)
  - 
  -

# Implementing Locks

- Mutex Lock: a synchronization primitive
- AcquireLock(L) 
  - Done before critical section of code
  - Returns when safe for process to enter critical section



- ReleaseLock(L) 
  - Done after critical section
  - Allows another process to acquire lock



# Implementing Locks

Process

```
int L=0;
```

```
AcquireLock(L):
```

```
while (L==1) ;  
L = 1;
```

wait till  
lock released

once  
lock is  
released } return

Spin lock

/\* BUSY WAITING \*/

# Implementing Locks

*lock variable.*

```
int L=0;
```

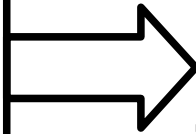
```
AcquireLock(L):  
    while (L==1) ;  
    L = 1;
```

```
ReleaseLock(L):  
    L = 0;
```

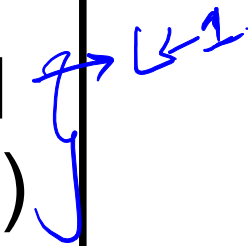
*/\* BUSY WAITING \*/*

# Problem in Implementing Locks

```
AcquireLock(L):  
  while (L==1) ;  
  L = 1;
```



```
wait: LW R1, Addr(L)  
      BNEZ wait  
      ADDI R1, R1, 1  
      SW R1, Addr(L)
```



# Problem in Implementing Locks

```
wait: LW R1, Addr(L)
      BNEZ wait
      ADDI R1, R1, 1
      SW R1, Addr(L)
```

Initially  $L=0$ . P1 and P2 are in contention to acquire the lock.

# Problem in Implementing Locks

**Process 1**

**Process 2**

```
wait: LW R1, Addr(L)
      BNEZ wait
      ADDI R1, R1, 1
      SW R1, Addr(L)
```

Initially  $L=0$ . P1 and P2 are in contention to acquire the lock.



# Problem in Implementing Locks

**Process 1**

**Process 2**

*L=0*  
**LW R1, Addr(L)**

**Context Switch**

```
wait: LW R1, Addr(L)
      BNEZ wait
      ADDI R1, R1, 1
      SW R1, Addr(L)
```

Initially L=0. P1 and P2 are in contention to acquire the lock.



# Problem in Implementing Locks

**Process 1**

**LW R1, Addr(L)**

*R1=0*

**Context Switch**

**Process 2**

**LW R1, Addr(L)**

**BNEZ wait** *← R1=0*

**ADDI R1, R1, 1** *←*

**# Critical Section #** *←*

**Context Switch**



```
wait: LW R1, Addr(L)
      BNEZ wait
      ADDI R1, R1, 1
      SW R1, Addr(L)
```

Initially L=0. P1 and P2 are in contention to acquire the lock.

# Problem in Implementing Locks

**Process 1**

**LW R1, Addr(L)**

**Context Switch**

**Process 2**

**LW R1, Addr(L)**

**BNEZ wait**

**ADDI R1, R1, 1**

**# Critical Section #**

**Context Switch**

**BNEZ wait**

**ADDI R1, R1, 1**

**# Critical Section #**

```
wait: LW R1, Addr(L)
      BNEZ wait
      ADDI R1, R1, 1
      SW R1, Addr(L)
```

Initially L=0. P1 and P2 are in contention to acquire the lock.

# Problem in Implementing Locks

**Process 1**

**LW R1, Addr(L)**

**Context Switch**

**BNEZ wait**

**ADDI R1, R1, 1**

**# Critical Section #**

**Process 2**

**LW R1, Addr(L)**

**BNEZ wait**

**ADDI R1, R1, 1**

**# Critical Section #**

**Context Switch**

wait: LW R1, Addr(L)  
BNEZ wait  
ADDI R1, R1, 1  
SW R1, Addr(L)

Initially L=0. P1 and P2 are in contention to acquire the lock.

**Both P1 and P2 are  
executing in the  
Critical Section !!!**

- indivisible
- all or none

# Atomic Exchange

- Hardware support for lock implementation

# Atomic Exchange

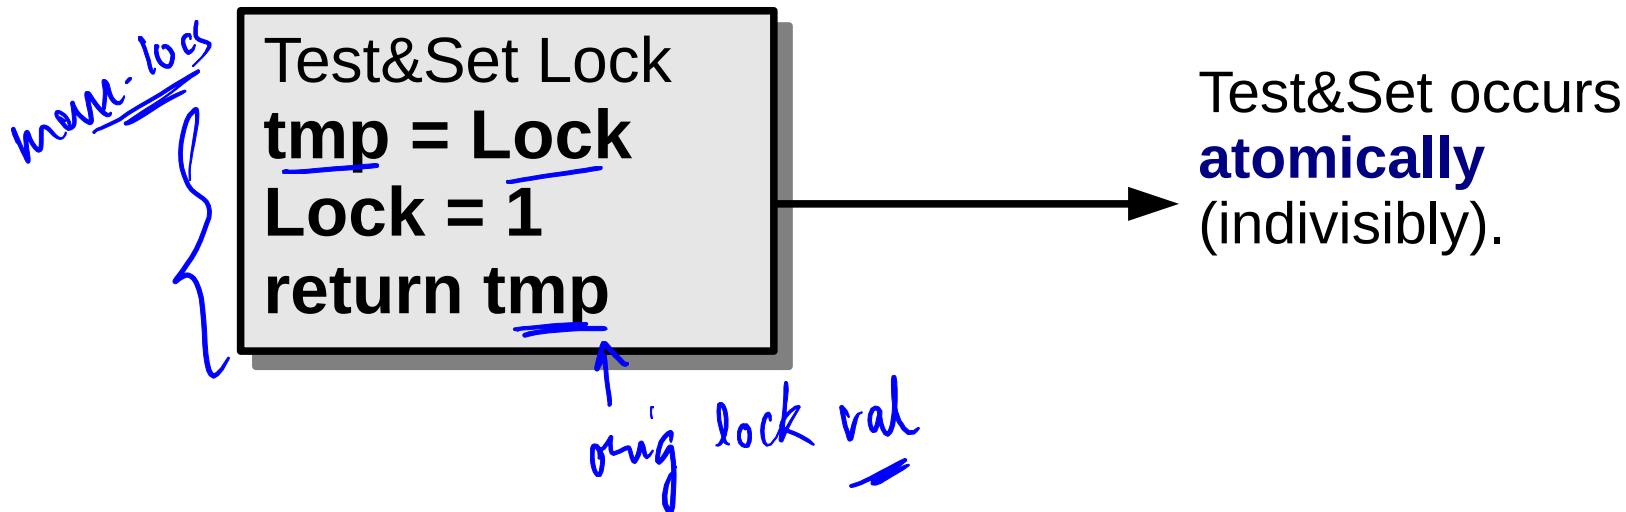
- Hardware support for lock implementation
- Atomic exchange: Swap contents between register and memory.

# Atomic Exchange

- Test&Set

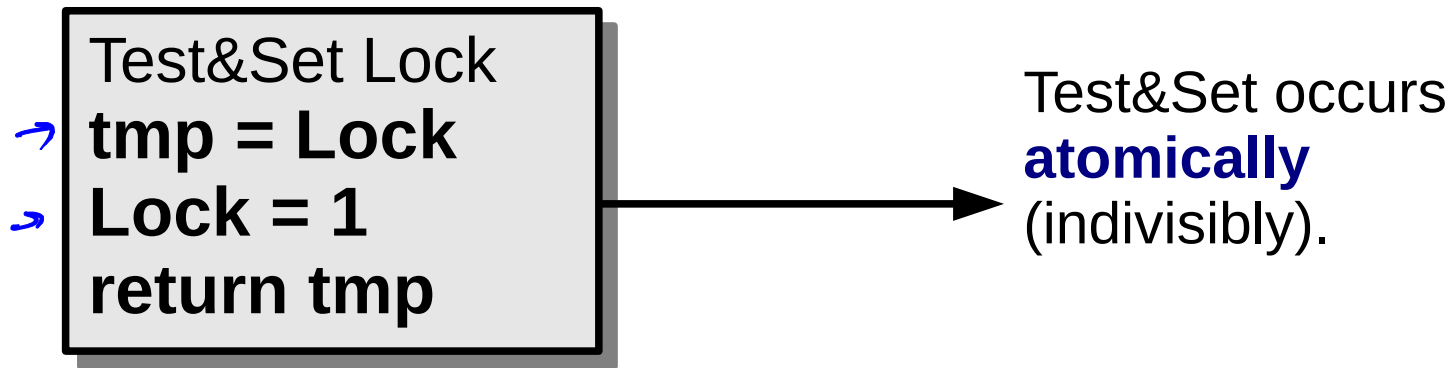
*test of lock = 0*  
*set lock = 1*

- Takes one memory operand and a register operand



# Atomic Exchange

- Test&Set
  - Takes one memory operand and a register operand



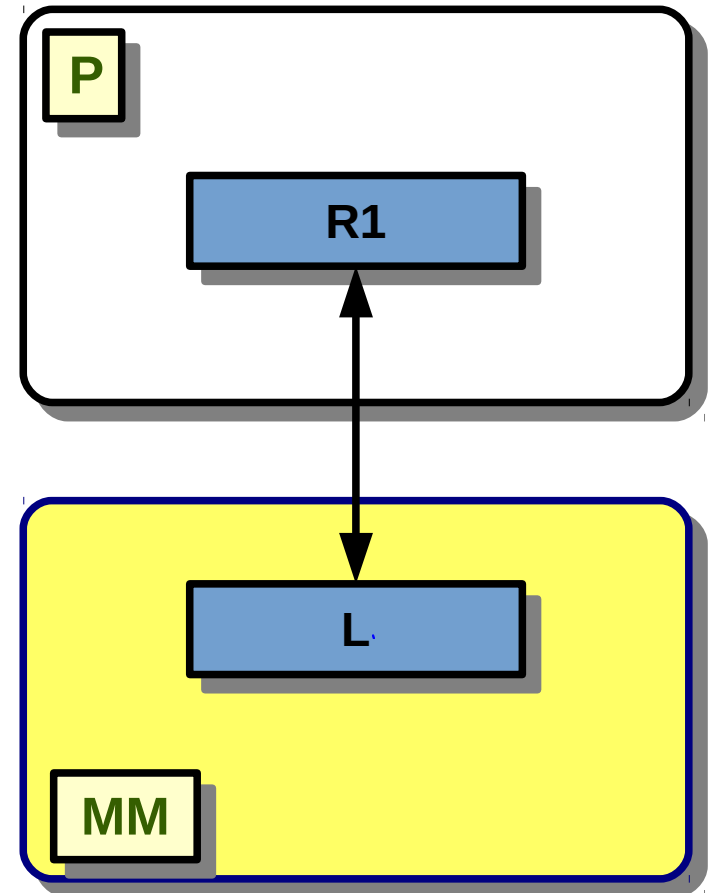
Atomic **Read-Modify-Write (RMW)** instruction

# Lock Implementation

*R1 → 1 #acquired*

*R1 has the orig L value.*

```
lock: Test&Set R1, L
      BNZ R1, lock
      Critical Section
      SW R0, L
```

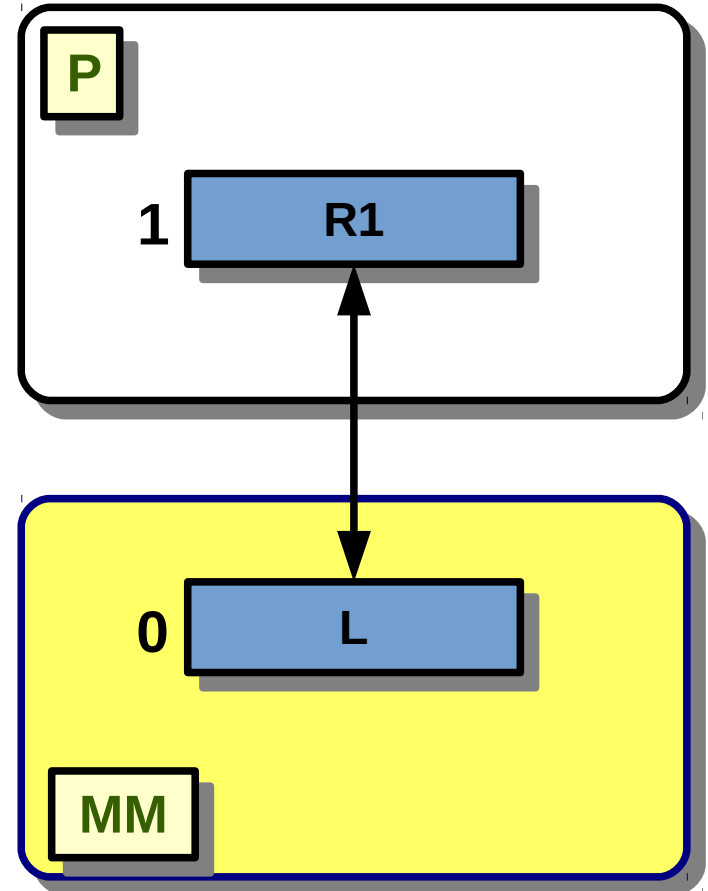




# Lock Implementation

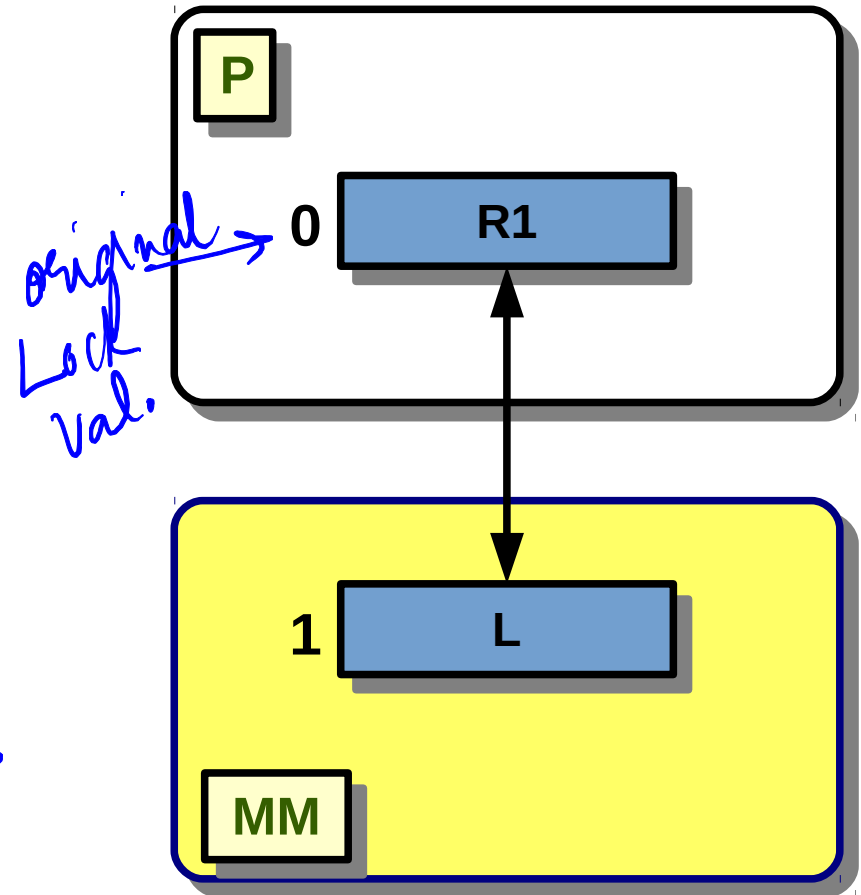
lock: Test&Set R1, L  
BNZ R1, lock  
Critical Section  
SW R0, L

① before T&S



# Lock Implementation

lock: Test&Set R1, L  
BNZ R1, lock  
Critical Section  
SW R0, L

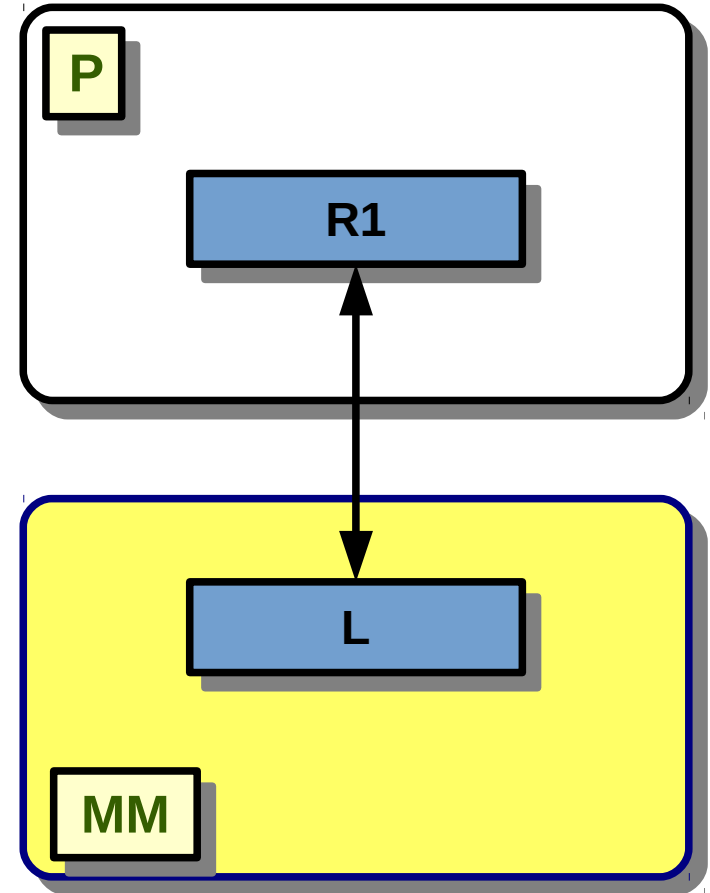


# Test&Set Implementation

Test&Set R1, L

```
t = L;   # Store a copy of Lock  
L = 1;   # Set Lock  
R1 ← t;  # Write original value  
         of the Lock in R1
```

{ stores: L, tmp  
 loads: R1

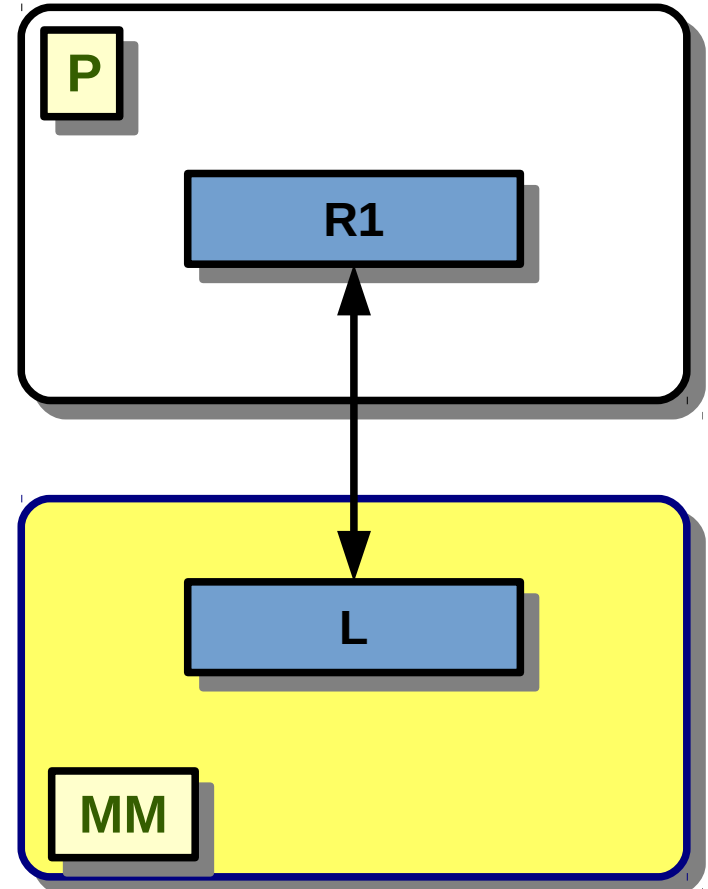


# Test&Set Implementation

**Test&Set R1, L**

**t = L;    # Store a copy of Lock  
L = 1;    # Set Lock  
R1 ← t;   # Write original value  
          of the Lock in R1**

**2 stores (t, L) and 1 load (R1)  
(atomic – practically 1  
instruction)**



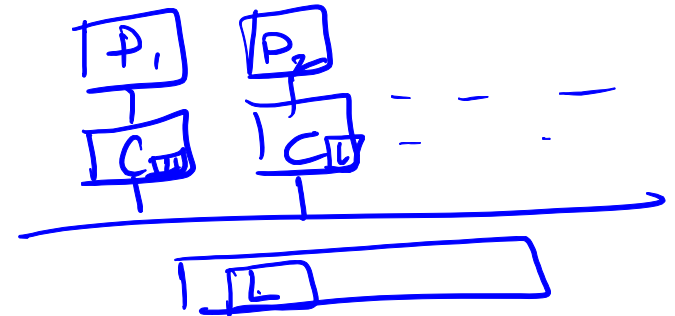
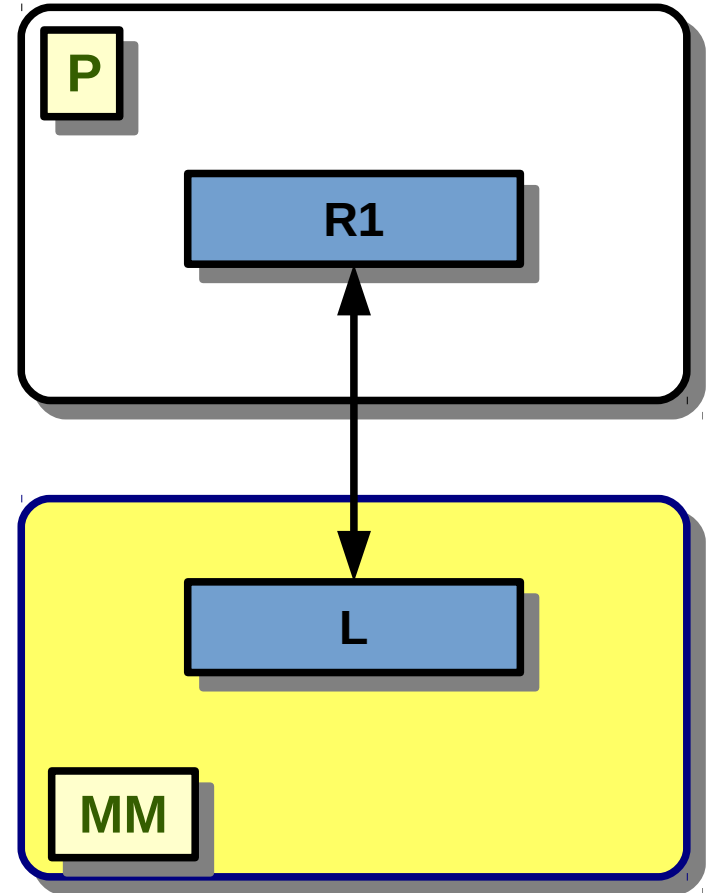
# Test&Set Implementation

Test&Set R1, L

$t = L;$  # Store a copy of Lock  
 $L = 1;$  # Set Lock  
 $R1 \leftarrow t;$  # Write original value  
of the Lock in R1

2 stores (t, L) and 1 load (R1)  
(atomic – practically 1  
instruction)

The atomic **read-modify-write** hardware  
primitive facilitates synchronization  
implementations (locks, barriers, etc.)

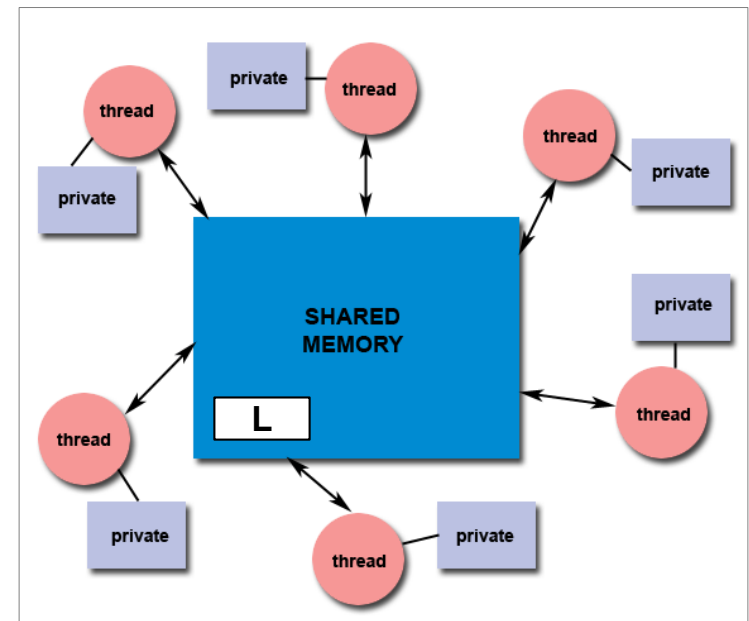


# Test&Set – Issues

- T&S is an atomic Read-Modify-Write instruction
  - 2 Stores, 1 Load
  - Atomic

# Test&Set – Issues

- T&S is an atomic Read-Modify-Write instruction
  - 2 Stores, 1 Load
  - Atomic
- Consider N processors executing T&S on a single lock variable



# Test&Set – Issues

- T&S is an atomic Read-Modify-Write instruction
  - 2 Stores, 1 Load
  - Atomic
- Consider N processors executing T&S on a single lock variable



# Test&Set – Issues

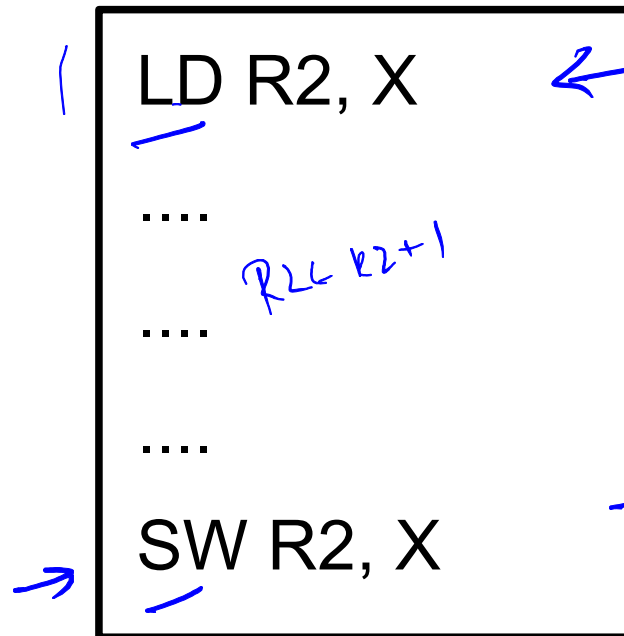
- T&S is an atomic Read-Modify-Write instruction
  - 2 Stores, 1 Load
  - Atomic
- Consider N processors executing T&S on a single lock variable
  - Redundant, useless work
  - Performance loss

# Test&Set – Issues

- T&S is an atomic Read-Modify-Write instruction
  - 2 Stores, 1 Load
  - Atomic
- Use LL-SC



# LL-SC Example



← read lock

→ update lock

$t_1$ :  $P_1$  has lock  
 $P_2 \dots P_n$ : spin wait.

$t_2$ :  $P_1$  releases lock

$P_2 \dots P_n$  } attempt acquire lock

$t_2 + \delta$ :  $P_2$  wins

$P_3 \dots P_n \rightarrow$  wasted work

# LL-SC Example

Atomic execution required!

LD R2, X

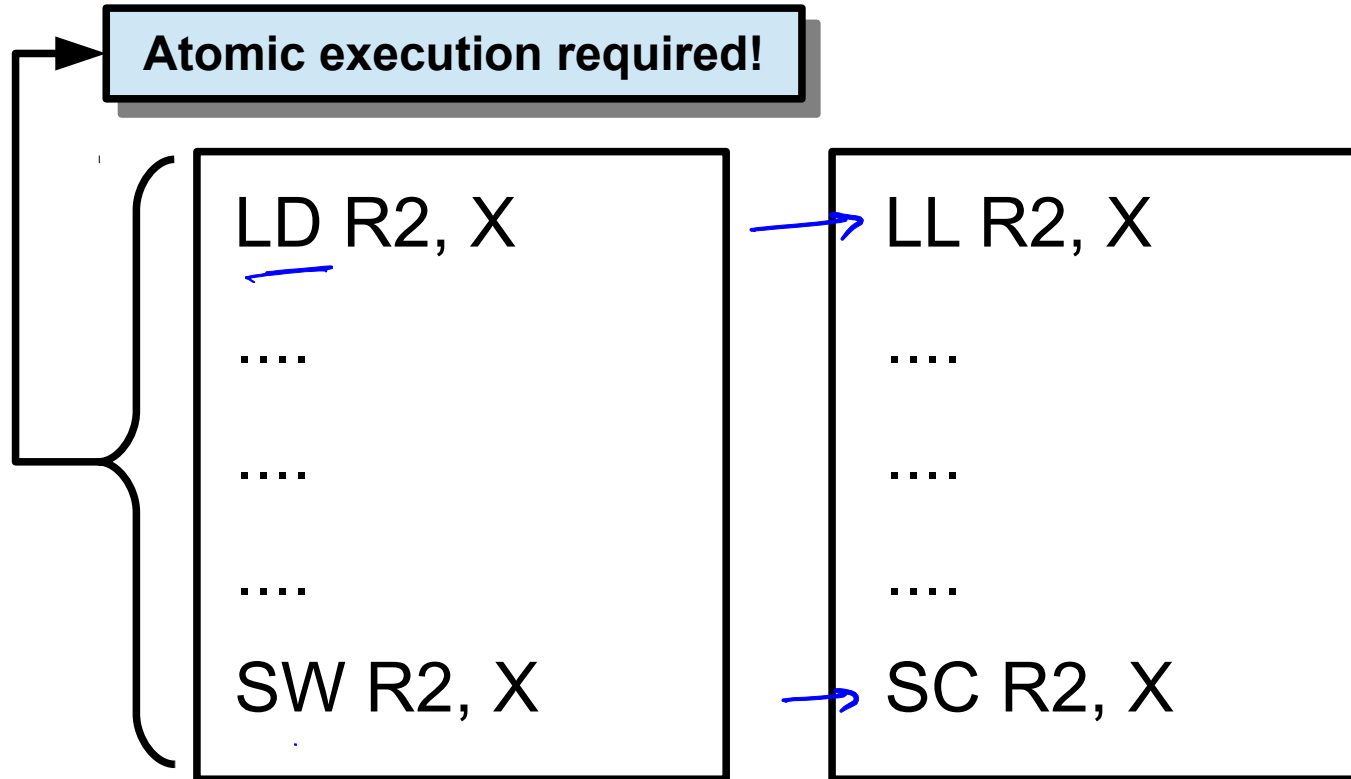
....

....

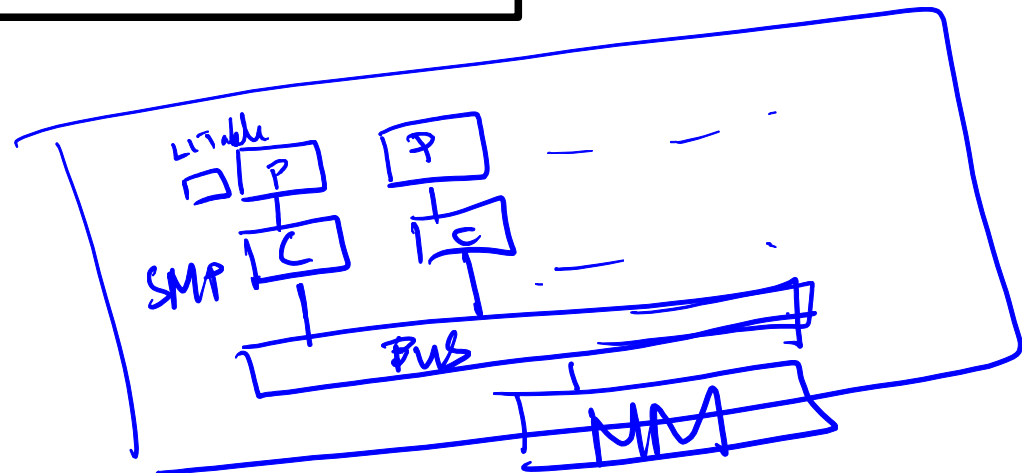
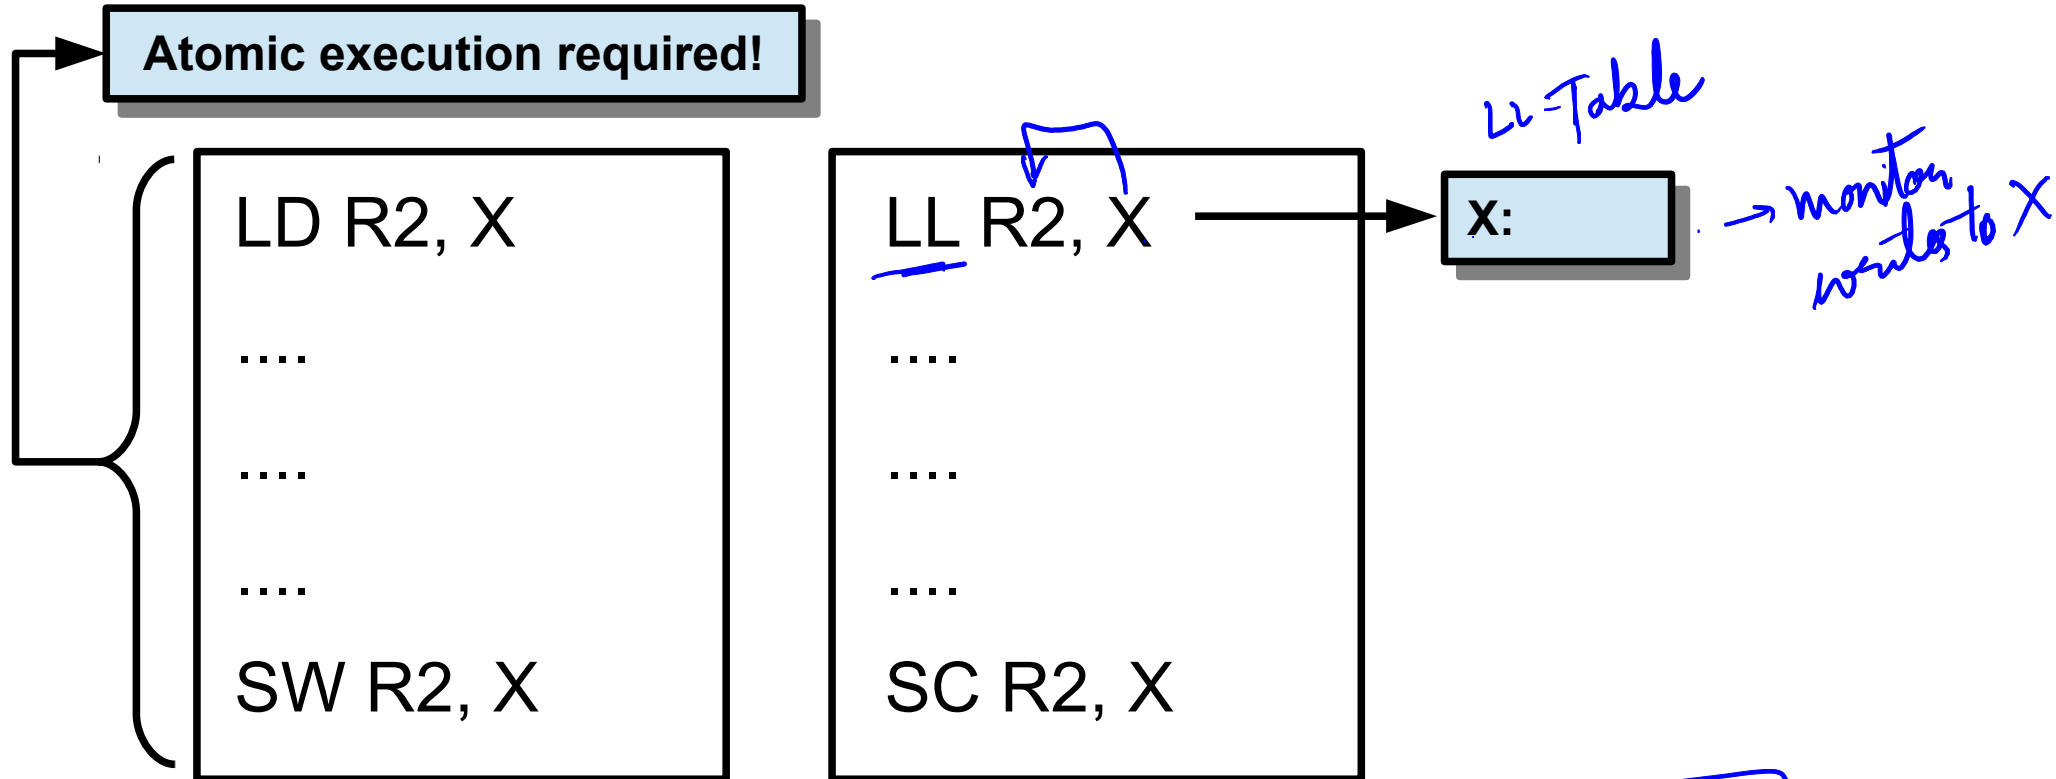
....

SW R2, X

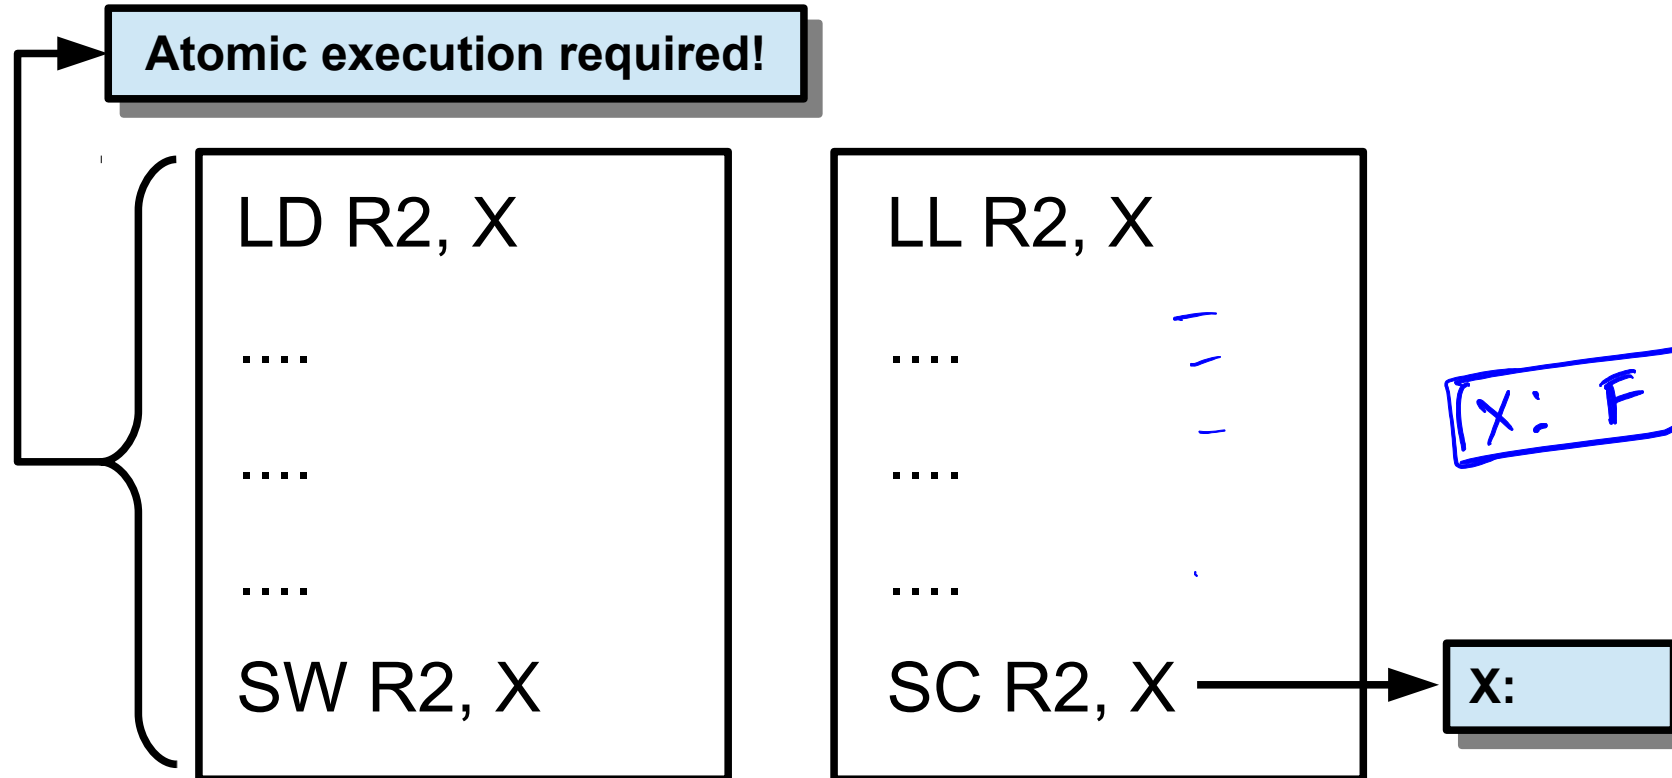
# LL-SC Example



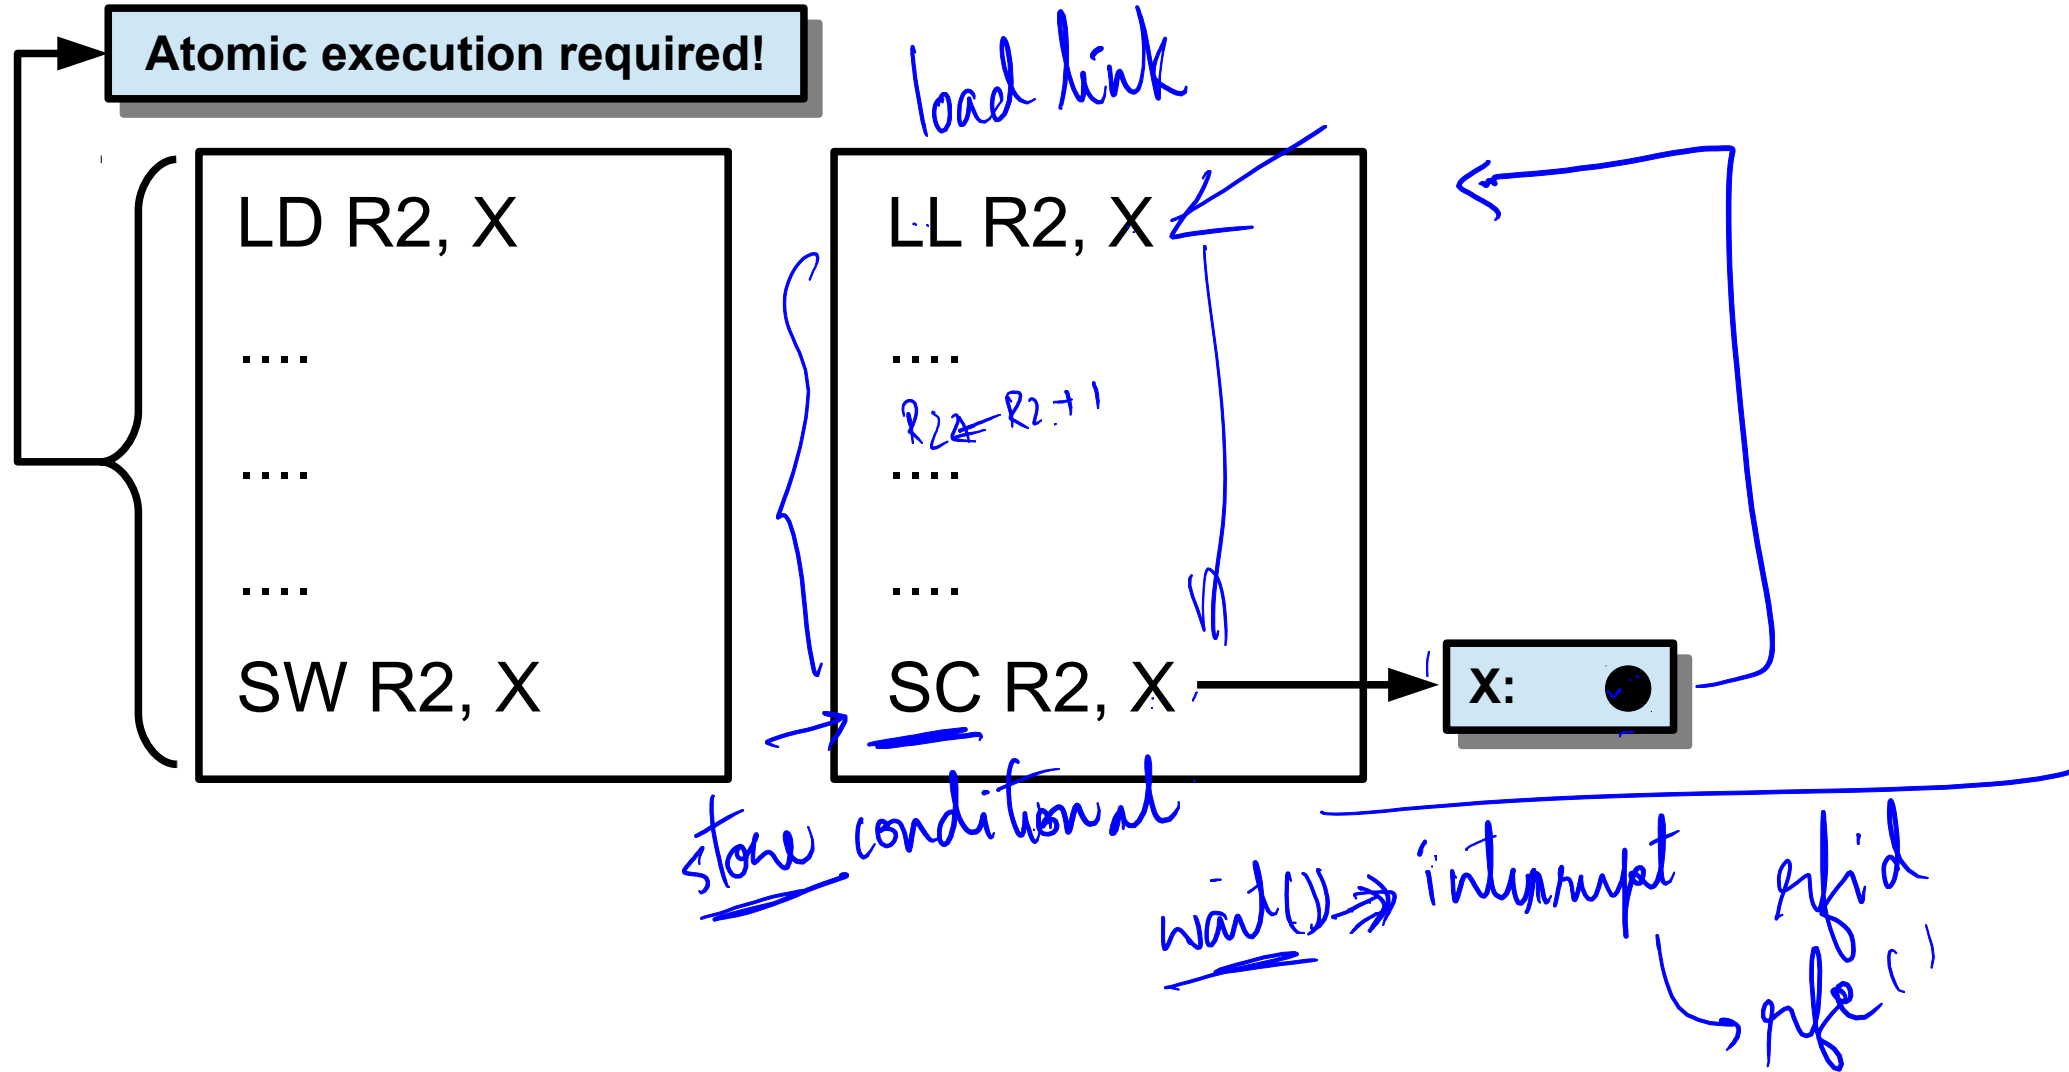
# LL-SC Example



# LL-SC Example

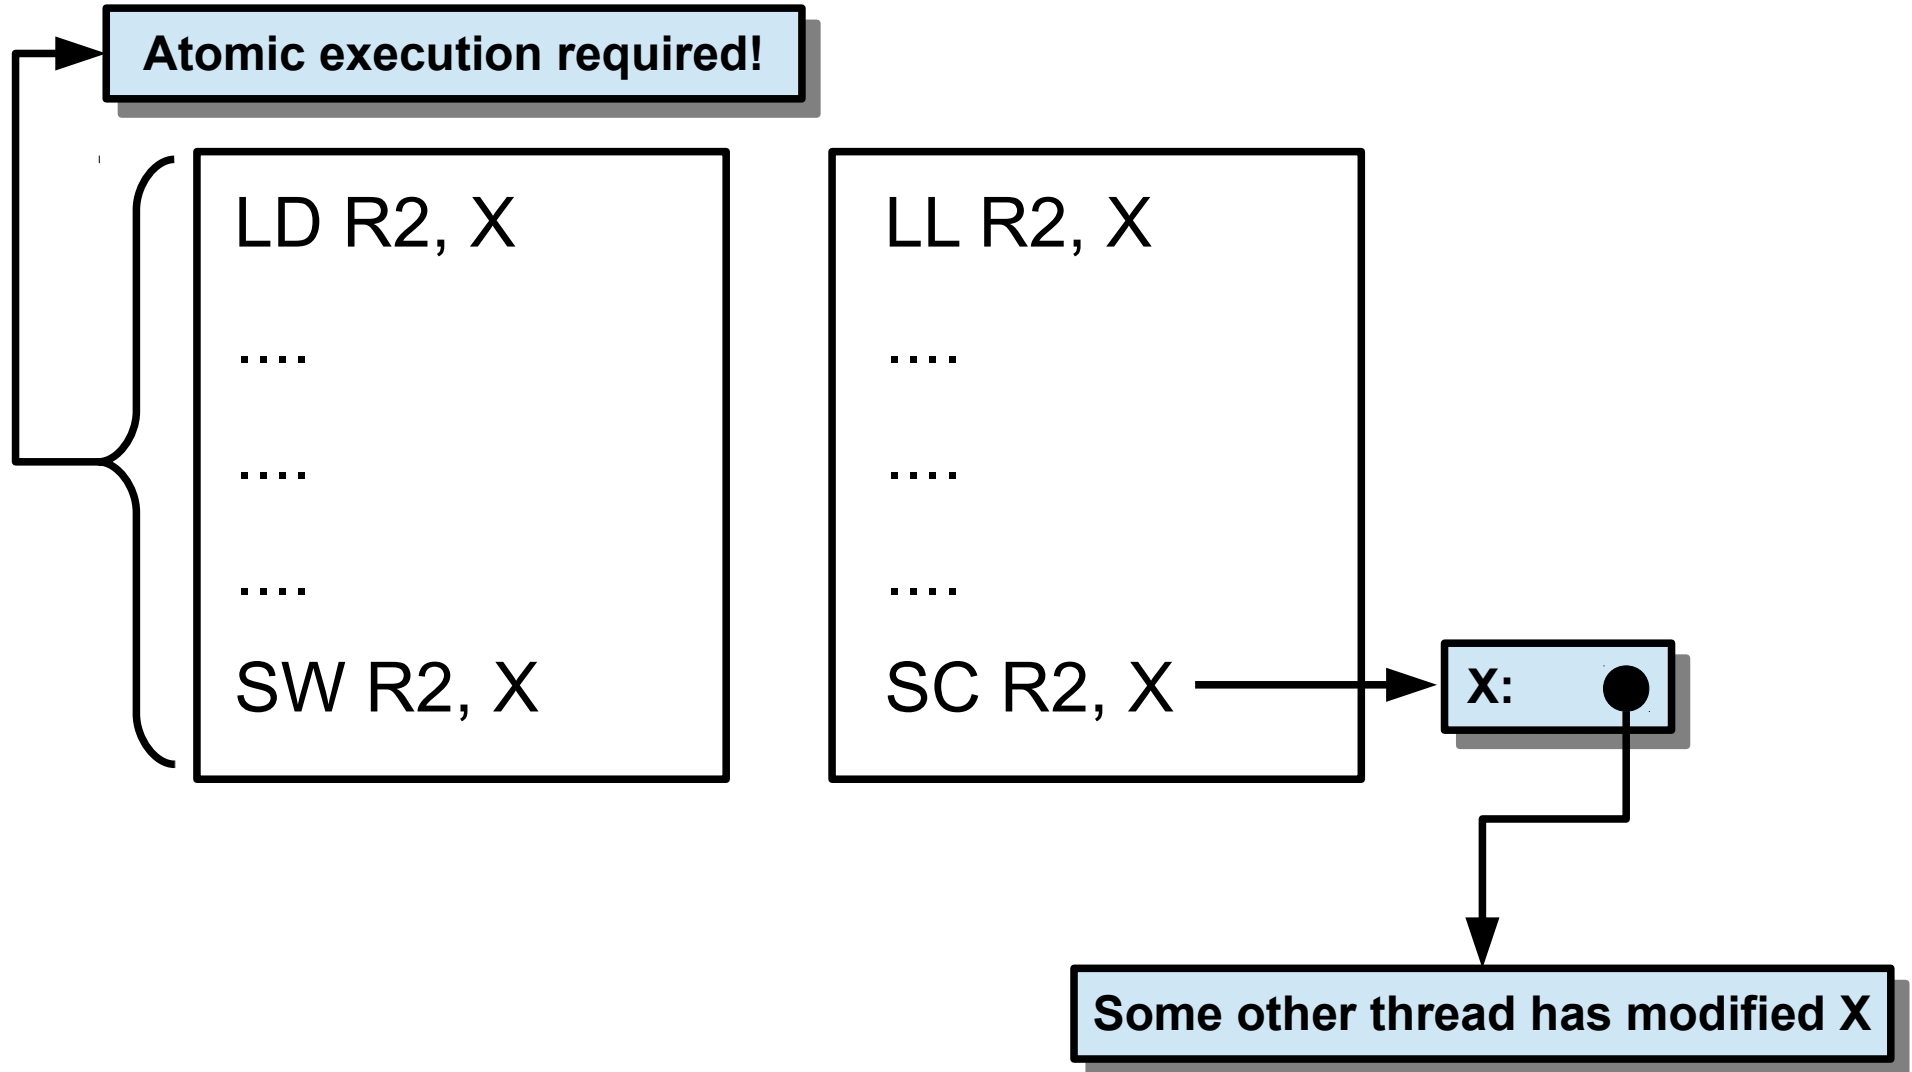


# LL-SC Example

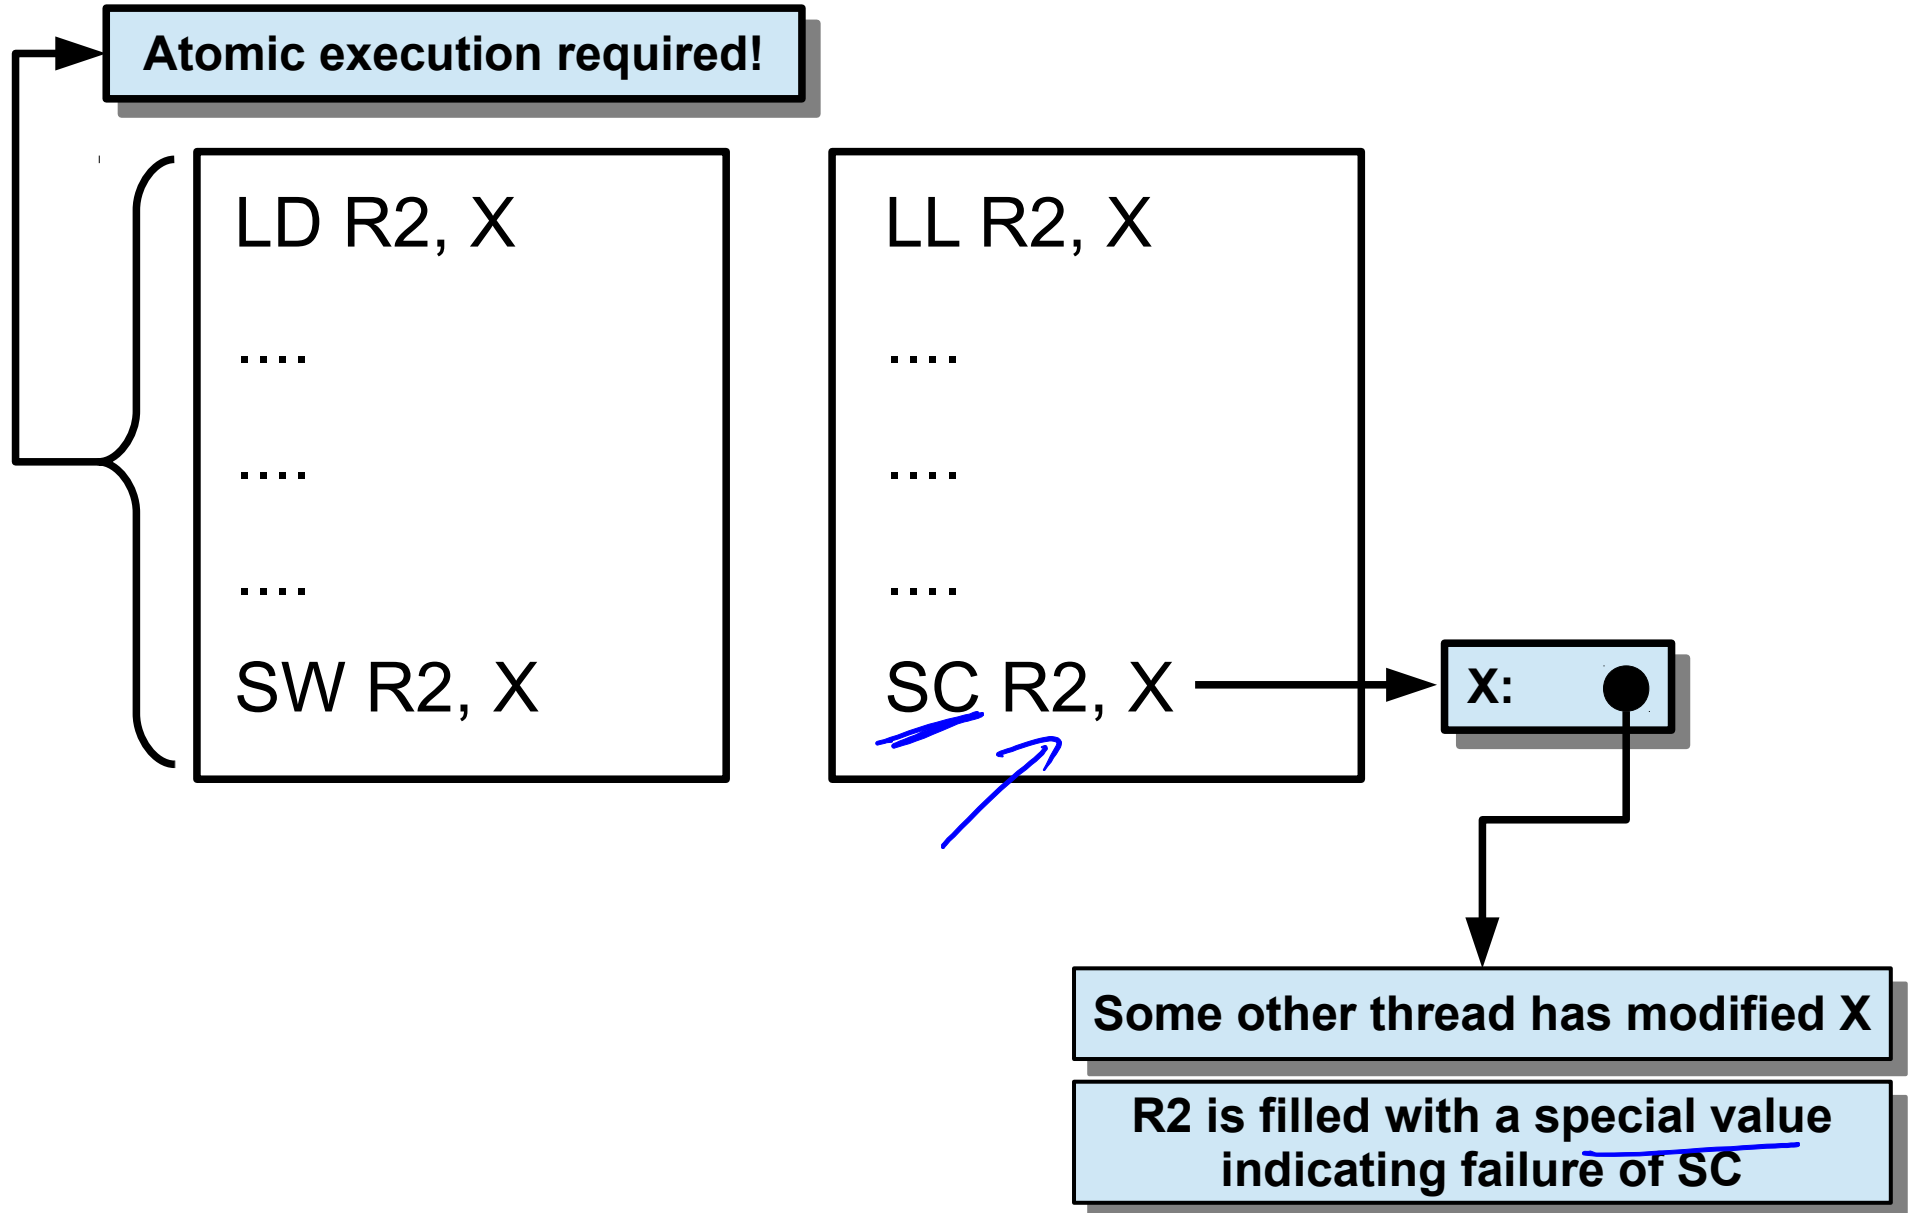




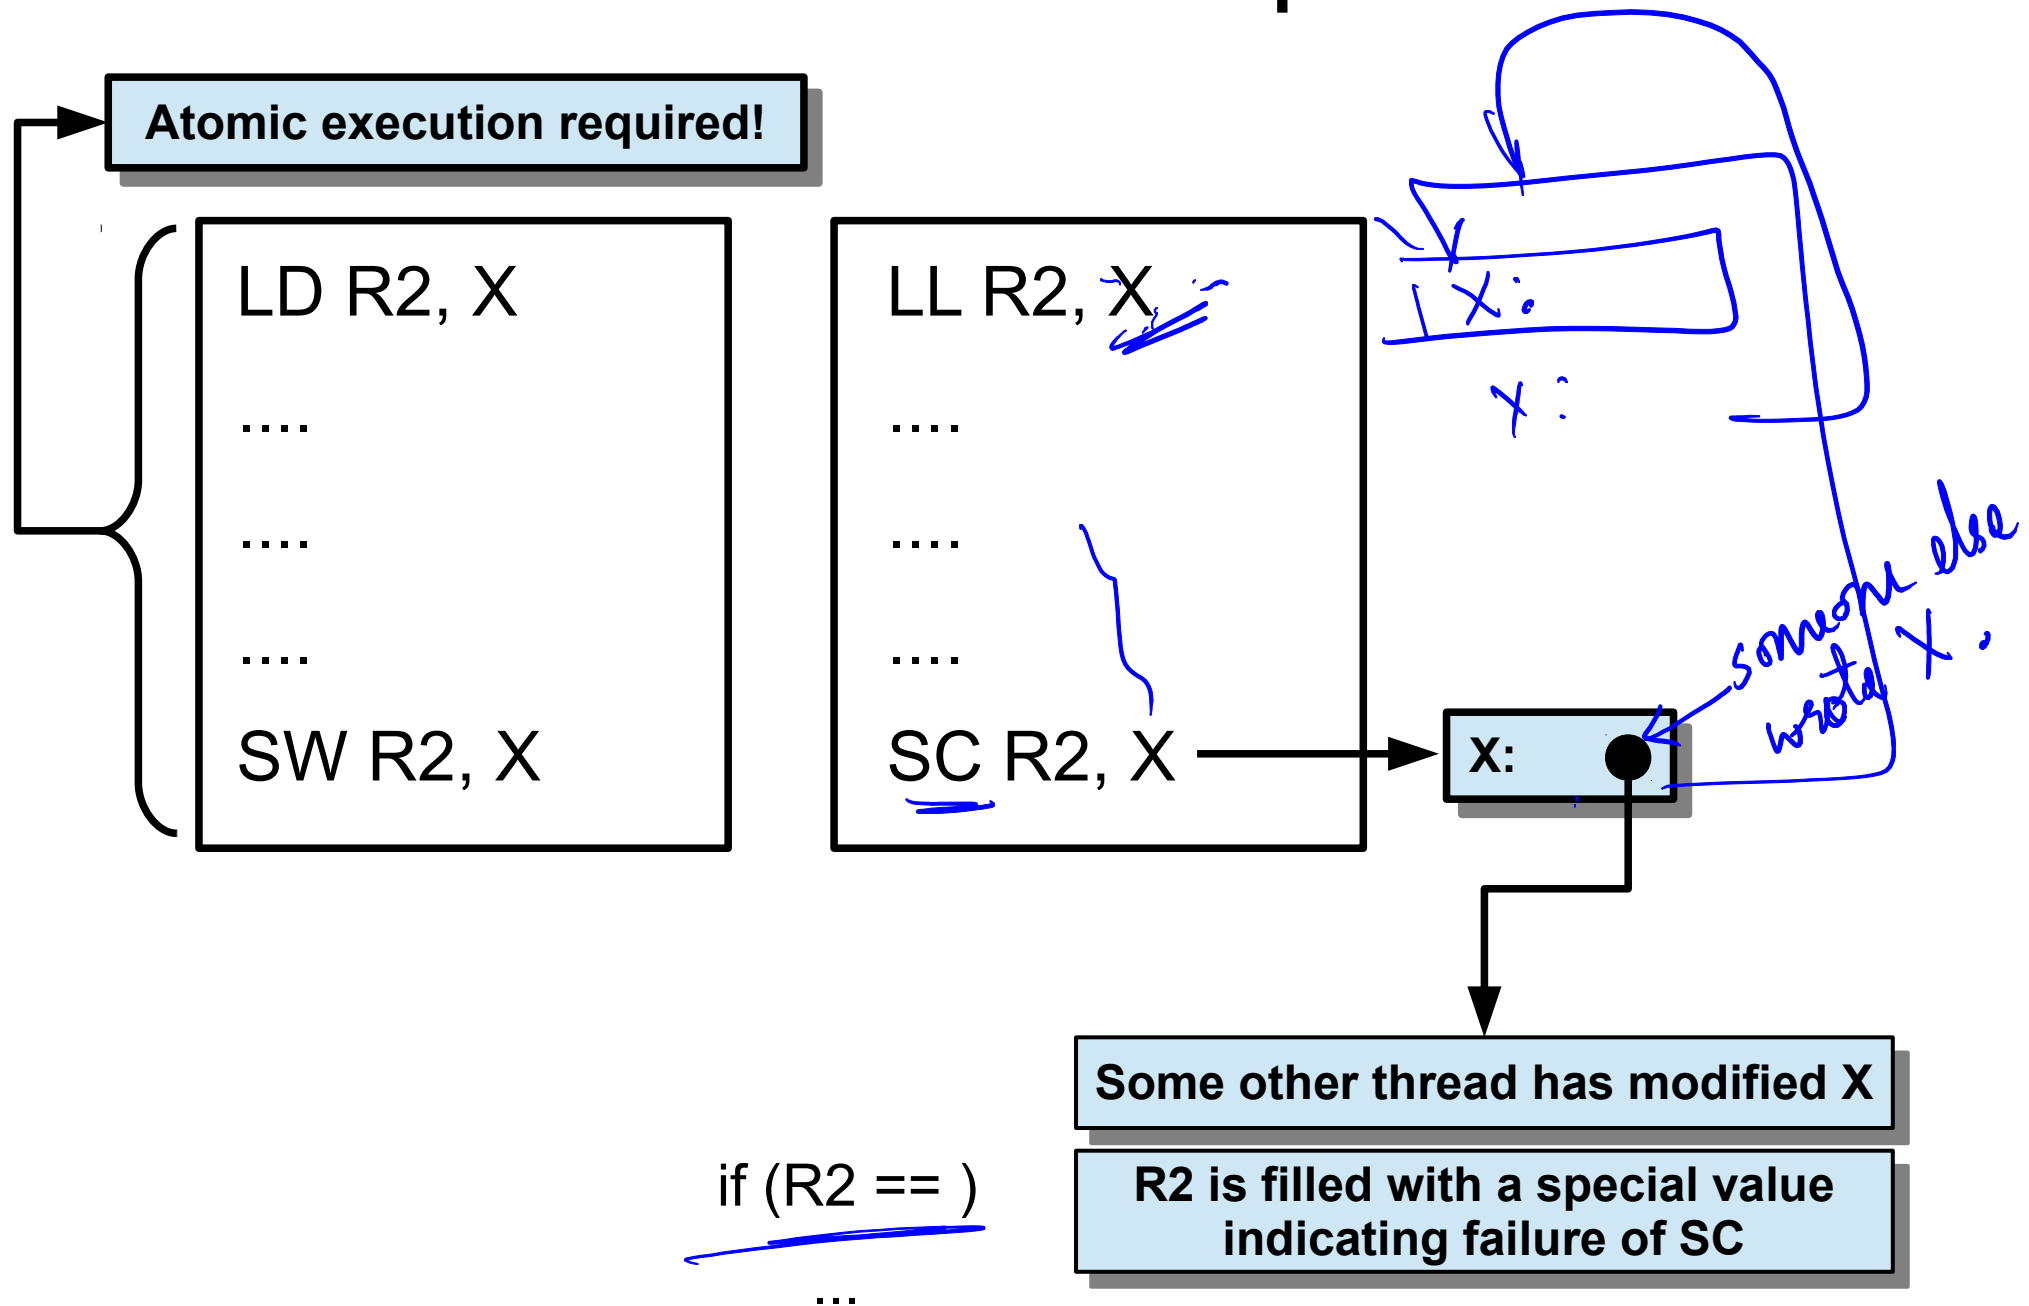
# LL-SC Example



# LL-SC Example



# LL-SC Example



# Load Linked and Store Conditional

- LL records the lock variable address in a table

# Load Linked and Store Conditional

- LL records the lock variable address in a table
- Snoop on the bus; Record the first Write Invalidate request on the lock variable in the table.

# Load Linked and Store Conditional

- LL records the lock variable address in a table
- Snoop on the bus; Record the first Write Invalidate request on the lock variable in the table.
- Before the thread updates its cache copy, SC checks the table.

# Load Linked and Store Conditional

- LL records the lock variable address in a table
- Snoop on the bus; Record the first Write Invalidate request on the lock variable in the table.
- Before the thread updates its cache copy, SC checks the table.
- If flag is set, SC fails (no coherence traffic). Fail is indicated by a special value in the register.

# Load Linked and Store Conditional

- LL records the lock variable address in a table
- Snoop on the bus; Record the first Write Invalidate request on the lock variable in the table.
- Before the thread updates its cache copy, SC checks the table.
- If flag is set, SC fails (no coherence traffic). Fail is indicated by a special value in the register.
- If flag is not set, SC succeeds. Lock acquired.



# Spin Lock using LL-SC

lockit: LL R2, 0(R1) ; no coherence traffic  
BNEZ R2, lockit ; not available, keep spinning  
DADDUI R2, R0, #1 ; put value 1 in R2  
SC R2, 0(R1) ; store-conditional succeeds if no one  
; updated the lock since the last LL  
BEQZ R2, lockit ; confirm that SC succeeded, else keep trying

*Transactional  
Memory*

# Spin Lock using LL-SC

Spin lock with  
lower coherence  
traffic.

lockit:	LL R2, 0(R1)	; no coherence traffic
	BNEZ R2, lockit	; not available, keep spinning
	DADDUI R2, R0, #1	; put value 1 in R2
	SC R2, 0(R1)	; store-conditional succeeds if no one ; updated the lock since the last LL
	BEQZ R2, lockit	; confirm that SC succeeded, else keep trying

# Outline

- Parallelism
- Flynn's classification
- Vector Processing
  - Subword Parallelism
- Symmetric Multiprocessors, Distributed Memory Machines
  - Shared Memory Multiprocessing, Message Passing
- Synchronization Primitives
  - Locks, LL-SC
- Cache coherence